

Open-LLM-V-Tuber

In diesem kleinen Projekt geht es um den Versuch den virtuellen, KI gesteuerten Chatbot „Open-LLM-V-Tuber“ zum Laufen zu bekommen. Der Chatbot selbst stammt vom chinesischen Entwickler Yi-Ting Chiu

Der Container

Zuerst möchte ich versuchen, den Chatbot in einem Container zu hosten. Dazu erstelle ich mir folgenden Container

Create: LXC Container

General

Template

Disks

CPU

Memory

Network

DNS

Confirm

Key ↑	Value
cores	4
hostname	Hatsukaze
memory	2048
nameserver	192.168.2.7
net0	name=eth0,bridge=vbr0,firewall=1,ip6=dhcp,ip=dhcp
nodename	soryu
ostemplate	local:vztmpl/ubuntu-22.04-standard_22.04-1_amd64.tar.zst
pool	
rootfs	local:25
searchdomain	sotoba.de
ssh-public-keys	
swap	2048
vmid	102

☐ Start after created

Advanced ☒

Back

Finish

Nach dem Erstellen des Containers, bringe ich ihn erstmal auf den neusten Stand

```
apt-get update && apt-get upgrade && reboot
```

[illegible]

Nach dem Neustart installiere ich dann die nötigen Voraussetzungen für den Chatbot. Zuerst Git

```
apt install git
```

```
root@natsukaze:~# apt install git
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  git-man libbrotli1 libcurl3-gnutls liberror-perl libgdbm-compat4 libperl5.34 librtmp1 libssh-4 patch perl perl-modules-5.34
Suggested packages:
  git-daemon-run | git-daemon-sysvinit git-doc git-email git-gui gitk gitweb git-cvs git-mediawiki git-svn diffutils-doc perl-doc libterm-readline-gnu-perl
  | libterm-readline-perl-perl make libtap-harness-archive-perl
The following NEW packages will be installed:
  git git-man libbrotli1 libcurl3-gnutls liberror-perl libgdbm-compat4 libperl5.34 librtmp1 libssh-4 patch perl perl-modules-5.34
0 upgraded, 12 newly installed, 0 to remove and 5 not upgraded.
Need to get 13.1 MB of archives.
After this operation, 71.5 MB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Danach ffmpeg

```
apt install ffmpeg
```

[illegible]

Als nächstes brauche ich das Python Umgebungsmanagement

```
wget -qO- https://astral.sh/uv/install.sh | sh
```

```
root@htcsukars:~# wget -qO- https://astral.sh/uv/install.sh | sh
downloading uv 0.7.8 x86_64-unknown-linux-gnu
--2025-05-27 16:36:33-- https://github.com/astral-sh/uv/releases/download/0.7.8/uv-x86_64-unknown-linux-gnu.tar.gz
Resolving github.com (github.com)... 140.82.121.3
Connecting to github.com (github.com)[140.82.121.3]:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: /objects/githubusercontent.com/github-production-release-asset-2e65be/699532645/55b0de0c-6b6e-4516-b6f2-50931b850e08?X-Amz-Algorithm=AWA3-4-IMAC-SHA256xK-AWZ-Credential=releases.asset.production%2F20250527%2Fus-east-1%2Fa3%2Paw4_request%3X-Amz-Date=20250527T16352824X-Amz-Expires=3004X-Amz-Signature=3e92c293128d493816559591-g76C8r7Y53feb38f7cb8167B9R631c6c484ec34c9b4X-Amz-SignedHeaders=host&response-content-disposition=attachment%3Bfilename%3Duv-x86_64-unknown-linux-gnu.tar.gz;response-content-type=application%2Foctet-stream [following]
--2025-05-27 16:36:33-- https://objects.githubusercontent.com/github-production-release-asset-2e65be/699532645/55b0de0c-6b6e-4516-b6f2-50931b850e08?X-Amz-Algorithm=AWA3-4-IMAC-SHA256xK-AWZ-Credential=releases.asset.production%2F20250527%2Fus-east-1%2Fa3%2Paw4_request%3X-Amz-Date=20250527T16352824X-Amz-Expires=3004X-Amz-Signature=3e92c293128d493816559591-g76C8r7Y53feb38f7cb8167B9R631c6c484ec34c9b4X-Amz-SignedHeaders=host&response-content-disposition=attachment%3Bfilename%3Duv-x86_64-unknown-linux-gnu.tar.gz;response-content-type=application%2Foctet-stream
Resolving objects.githubusercontent.com (objects.githubusercontent.com)... 185.199.111.133, 185.199.109.133, 185.199.108.133, ...
Connecting to objects.githubusercontent.com (objects.githubusercontent.com)[185.199.111.133]:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 17486328 (17M) [application/octet-stream]
Saving to: '/tmp/.KFnygFKqrE/input.tar.gz'

/tmp/.KFnygFKqrE/input.tar.gz          47k[=====]           ]   7.91M  1.97MB/s    eta 5s
```

Sollte es benötigt werden, installiere ich auch curl.

apt install curl

```
root@Hatsukaze:~# apt install curl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  libcurl4
The following NEW packages will be installed:
  curl libcurl4
0 upgraded, 2 newly installed, 0 to remove and 5 not upgraded.
Need to get 483 kB of archives.
After this operation, 1262 kB of additional disk space will be used.
Do you want to continue? [Y/n] ☐
```

Active Directory

Da dieser Container Teil des Active Directory sein soll, bereite ich die Einladung vor.
Dazu installiere ich zuerst Samba.

```
apt install sssd-ad sssd-tools adcli realmd krb5-user libnss-sss libpam-sss  
samba-common-bin oddjob oddjob-mkhomedir
```

```
root@hatsukaze:~# apt install sssd-ad sssd-tools adcli realmd krb5-user libnss-sss libpam-sss samba-common-bin oddjob oddjob-mkhomedir
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  cracklib-runtime dirmngr gnupg gnupg-l10n gnupg-utils gpg gpg-agent gpg-wks-client gpg-wks-server gpgconf gpgsm krb5-config ldap-utils libassuan0 libavahi-client3
  libavahi-common-data libavahi-common3 libbasicobjects0 libc-ares2 libcollection4 libcrack2 libcups2 libdhash1 libgpgme11 libgsasl4 libini-config5 libipa-hbac0
  libkadm5clnt-mit12 libkadm5srv-mit12 libkdb5-10 libkksba8 libldap-2.5-0 libldap-common libldb2 libnfsidmap1 libnl-3-200 libnl-route-3-200 libnpt0 libpam-pwquality
  libpath-util1 libpolkit-agent-1-0 libpolkit-gobject-1-0 libpwquality-common libpwquality1 libpython3.10 libref-array1 libsass2-modules libsass2-modules-gssapi-mit
  libssmclient libsss-certmap0 libsss-idmap0 libsss-nss-idmap0 libtalloc2 libtdb1 libtevent0 libwbclient0 pinentry-curses pkexec policykit-1 polkitd python3-gpg
  python3-ldb python3-samba python3-sss python3-systemd python3-talloc python3-tdb samba-common samba-dsdb-modules samba-libe sssd sssd-ad-common sssd-common
  sssd-dbus sssd-ipa sssd-krb5 sssd-krb5-common sssd-ldap sssd-proxy wamerican
Suggested packages:
  dbus-user-session pinentry-gnome3 tor parcimonie xloadimage scdcaemon krb5-k5tls cups-common krb5-doc libsass2-modules-ldap libsass2-modules-otp libsass2-modules-sql
  pinentry-doc heimdal-clients python3-markdown python3-dnspython libsss-sudo
The following NEW packages will be installed:
  adcli cracklib-runtime dirmngr gnupg gnupg-l10n gnupg-utils gpg gpg-agent gpg-wks-client gpg-wks-server gpgconf gpgsm krb5-config krb5-user ldap-utils libassuan0
  libavahi-client3 libavahi-common-data libavahi-common3 libbasicobjects0 libc-ares2 libcollection4 libcrack2 libcups2 libdhash1 libgpgme11 libgsasl4 libini-config5
  libipa-hbac0 libkadm5clnt-mit12 libkadm5srv-mit12 libkdb5-10 libkksba8 libldap-2.5-0 libldap-common libldb2 libnfsidmap1 libnl-3-200 libnl-route-3-200 libnpt0
  libnss-sss libpam-pwquality libpam-sss libpath-util1 libpolkit-agent-1-0 libpolkit-gobject-1-0 libpwquality-common libpwquality1 libpython3.10 libref-array1
  libsass2-modules libsass2-modules-gssapi-mit libssmclient libsss-certmap0 libsss-idmap0 libsss-nss-idmap0 libtalloc2 libtdb1 libtevent0 libwbclient0 oddjob
  oddjob-mkhomedir pinentry-curses pkexec policykit-1 polkitd python3-gpg python3-ldb python3-samba python3-sss python3-systemd python3-talloc python3-tdb realmd
  samba-common samba-common-bin samba-dsdb-modules samba-libe sssd sssd-ad sssd-ad-common sssd-common sssd-dbus sssd-ipa sssd-krb5 sssd-krb5-common sssd-ldap
  sssd-proxy wamerican
0 upgraded, 90 newly installed, 0 to remove and 5 not upgraded.
Need to get 20.3 MB of archives.
After this operation, 85.8 MB of additional disk space will be used.
Do you want to continue? [Y/n]
```

Erstelle die sssd.conf Datei mit folgenden Daten.

```
nano /etc/sss/sss.conf
```

```
[sss]
```

```
services = nss, pam, ssh
```

```
config_file_version = 2
```

```
domains = sotoba.de
```

```
[domain/sotoba.de]
```

```
id_provider = ad
```

```
access_provider = ad
```

```
ad_domain = sotoba.de
```

```
krb5_realm = SOToba.DE
```

```
realmd_tags = manages-system joined-with-adcli
```

```
cache_credentials = True
```

```
default_shell = /bin/bash
```

```
fallback_homedir = /home/%u@%d
```

```
ldap_id_mapping = True
```

```
use_fully_qualified_names = False
```

Ändere die Berechtigungen der Datei

```
chmod 600 /etc/sss/sss.conf
```

```
chown root:root /etc/sss/sss.conf
```

```
ls -l /etc/sss/sss.conf
```

```
root@Hatsukaze:~# chmod 600 /etc/sss/sss.conf
chown root:root /etc/sss/sss.conf
ls -l /etc/sss/sss.conf
-rw----- 1 root root 365 May 27 16:18 /etc/sss/sss.conf
root@Hatsukaze:~#
```

Und lade den Container in die Domäne ein und Schau und nach einem reboot nach ob der Dienst richtig gestartet wurde

```
adcli join sotoba.de
```

```
reboot
```

```
systemctl status sssd
```

```
root@Hatsukaze:~# adcli join sotoba.de
Password for Administrator@SOToba.DE:
root@Hatsukaze:~# systemctl status sssd
* sssd.service - System Security Services Daemon
   Loaded: loaded (/lib/systemd/system/sss.service; enabled; vendor preset: enabled)
   Active: active (running) since Tue 2025-05-27 16:21:52 UTC; 385ms ago
     Main PID: 2309 (sss)
       Tasks: 5 (limit: 9329)
      Memory: 50.5M
         CPU: 1.517s
    CGroup: /system.slice/sss.service
            |-2309 /usr/sbin/sss -i --logger=files
            |-2319 /usr/libexec/sss/sss_be --domain sotoba.de --uid 0 --gid 0 --logger=files
            |-2321 /usr/libexec/sss/sss_nss --uid 0 --gid 0 --logger=files
            |-2322 /usr/libexec/sss/sss_ssh --uid 0 --gid 0 --logger=files
            \-2323 /usr/libexec/sss/sss_pam --uid 0 --gid 0 --logger=files

May 27 16:21:49 Hatsukaze sssd_ssh[2315]: Starting up
May 27 16:21:49 Hatsukaze sssd_pam[2314]: Starting up
May 27 16:21:49 Hatsukaze sssd_nss[2316]: Starting up
May 27 16:21:49 Hatsukaze sssd_ssh[2317]: Starting up
May 27 16:21:49 Hatsukaze sssd_pam[2318]: Starting up
May 27 16:21:51 Hatsukaze sssd_be[2319]: Starting up
May 27 16:21:51 Hatsukaze sssd_ssh[2322]: Starting up
May 27 16:21:52 Hatsukaze sssd_nss[2321]: Starting up
May 27 16:21:52 Hatsukaze sssd_pam[2323]: Starting up
May 27 16:21:52 Hatsukaze systemd[1]: Started System Security Services Daemon.
root@Hatsukaze:~#
```

Der Chatbot

Nun lade ich den Bot von Github runter

git clone https://github.com/Open-LLM-VTuber/Open-LLM-VTuber --recursive

```
root@Hatsukaze:~# git clone https://github.com/Open-LLM-VTuber/Open-LLM-VTuber --recursive
Cloning into 'Open-LLM-VTuber'...
remote: Enumerating objects: 5377, done.
remote: Counting objects: 100% (340/340), done.
remote: Compressing objects: 100% (106/106), done.
remote: Total 5377 (delta 271), reused 238 (delta 233), pack-reused 5037 (from 2)
Receiving objects: 100% (5377/5377), 45.46 MiB | 1.99 MiB/s, done.
Resolving deltas: 100% (3422/3422), done.
Submodule 'frontend' (https://github.com/Open-LLM-VTuber/Open-LLM-VTuber-Web) registered for path 'frontend'
Cloning into '/root/Open-LLM-VTuber/frontend'...
remote: Enumerating objects: 2954, done.
remote: Counting objects: 100% (711/711), done.
remote: Compressing objects: 100% (179/179), done.
remote: Total 2954 (delta 552), reused 581 (delta 532), pack-reused 2243 (from 1)
Receiving objects: 100% (2954/2954), 32.69 MiB | 2.01 MiB/s, done.
Resolving deltas: 100% (1852/1852), done.
remote: Enumerating objects: 5, done.
remote: Counting objects: 100% (3/3), done.
remote: Compressing objects: 100% (3/3), done.
remote: Total 5 (delta 0), reused 0 (delta 0), pack-reused 2 (from 1)
Unpacking objects: 100% (5/5), 574.25 KiB | 1.63 MiB/s, done.
From https://github.com/Open-LLM-VTuber/Open-LLM-VTuber-Web
 * branch          59808e04234378d50fc33d49275814191ab2a569 -> FETCH_HEAD
Submodule path 'frontend': checked out '59808e04234378d50fc33d49275814191ab2a569'
root@Hatsukaze:~#
```

Dann brauche ich das Programm uv, das ich mit wget heruntergeladen habe. Jedoch
zeit er bei dem Befehl

uv -version

```
root@Hatsukaze:~# uv --version
uv: command not found
```

Um das zu ändern, bearbeite ich die .bashrc Datei und füge folgende Anweisung ans
Ende der Datei

nano /root/.bashrc

export PATH=/root/.local/bin:\$PATH

```
GNU nano 6.2 /root/.bashrc
if [ -x /usr/bin/dircolors ]; then
    test -r ~/.dircolors && eval "$(dircolors -b ~/.dircolors)" || eval "$(dircolors -b)"
    alias ls='ls --color=auto'
    #alias dir='dir --color=auto'
    #alias vdir='vdir --color=auto'

    alias grep='grep --color=auto'
    alias fgrep='fgrep --color=auto'
    alias egrep='egrep --color=auto'
fi

# some more ls aliases
alias ll='ls -alF'
alias la='ls -A'
alias l='ls -CF'

# Alias definitions.
# You may want to put all your additions into a separate file like
# ~/.bash_aliases, instead of adding them here directly.
# See /usr/share/doc/bash-doc/examples in the bash-doc package.

if [ -f ~/.bash_aliases ]; then
    . ~/.bash_aliases
fi

# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile)
# sources /etc/bash.bashrc
if [ -f /etc/bash_completion ] && ! shopt -oq posix; then
    #
    . /etc/bash_completion
    #fi

. "$HOME/.local/bin/env"
export PATH=/root/.local/bin:$PATH
```

Ob es funktioniert hat. Zeigt sich nach einem reboot

```
Ubuntu 22.04.5 LTS Hatsukaze tty1

Hatsukaze login: root
Password:
Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.12-10-pve x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro
New release '24.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Tue May 27 16:28:09 UTC 2025 on lxc/tty1
root@Hatsukaze:~# uv --version
uv 0.7.8
root@Hatsukaze:~#
```

Nun Wechsel ich in das Stammverzeichnis des Bots und führe den Befehl uv sync aus, was den Chatbot installiert

```
root@Hatsukaze:~# cd Open-LLM-VTuber
root@Hatsukaze:~/Open-LLM-VTuber# uv sync
Using CPython 3.10.12 interpreter at: /usr/bin/python3.10
Creating virtual environment at: .venv
Resolved 275 packages in 9.27s
   Built srt==3.5.3
.: Preparing packages... (79/110)
protobuf ----- 959.49 KiB/1.00 MiB
pygments ----- 1007.45 KiB/1.17 MiB
soundfile ----- 992.00 KiB/1.25 MiB
aiohttp ----- 1020.59 KiB/1.52 MiB
networkx ----- 1.03 MiB/1.64 MiB
pydantic-core ----- 973.19 KiB/1.90 MiB
torchaudio ----- 939.00 KiB/3.22 MiB
uvloop ----- 948.81 KiB/3.65 MiB
virtualenv ----- 1004.79 KiB/4.10 MiB
silero-vad ----- 1022.80 KiB/4.79 MiB
sympy ----- 1.03 MiB/5.90 MiB
ruff ----- 994.38 KiB/10.76 MiB
onnxruntime ----- 990.05 KiB/12.71 MiB
nvidia-cuda-cupti-cu12 ----- 962.38 KiB/13.17 MiB
numpy ----- 991.18 KiB/17.40 MiB
nvidia-nvjitlink-cu12 ----- 1005.77 KiB/20.09 MiB
sherpa-onnx ----- 1010.74 KiB/21.16 MiB
nvidia-cuda-nvrtc-cu12 ----- 923.00 KiB/23.50 MiB
azure-cognitiveservices-speech ----- 930.34 KiB/37.89 MiB
scipy ----- 986.89 KiB/38.72 MiB
nvidia-cusolver-cu12 ----- 990.61 KiB/53.70 MiB
nvidia-cusparse-cu12 ----- 960.00 KiB/122.01 MiB
nvidia-cusparselt-cu12 ----- 963.56 KiB/143.11 MiB
nvidia-nccl-cu12 ----- 1005.03 KiB/179.91 MiB
nvidia-cusparse-cu12 ----- 1010.38 KiB/197.84 MiB
nvidia-cufft-cu12 ----- 974.67 KiB/201.66 MiB
triton ----- 981.62 KiB/241.37 MiB
nvidia-cublas-cu12 ----- 955.00 KiB/346.60 MiB
nvidia-cudnn-cu12 ----- 1005.08 KiB/633.96 MiB
torch ----- 1011.56 KiB/731.17 MiB
```

Nun starte ich das Hauptprogramm

uv run run_server.py

```
root@Hatsukaze:~/Open-LLM-VTuber# uv run run_server.py
2025-05-27 17:40:56.68 | INFO | __main__:modules:86 - Running in standard mode. For detailed debug logs, use: uv run run_server.py --verbose
2025-05-27 17:40:56 | INFO | __main__:run:37 | Open-LLM-VTuber, version v1.1.3
2025-05-27 17:40:56 | WARNING | upgrade:sync_user_config:352 | Warning: config.yaml not found
2025-05-27 17:40:56 | WARNING | upgrade:sync_user_config:353 | Copying default configuration from template
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.service_context:init_live2d:156 | Initializing Live2D: shizuku-local
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.live2d_model:lookup_model_info:142 | Model Information Loaded.
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.service_context:init_asr:166 | Initializing ASR: sherpa_onnx_asr
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.asr.sherpa_onnx_asr:_init_:61 | Sherpa-ONNX-ASR: Using cpu for inference
2025-05-27 17:40:56 | WARNING | src.open_llm_vtuber.asr.sherpa_onnx_asr:_create_recognizer:166 | SenseVoice model not found. Downloading the model...
2025-05-27 17:40:56 | WARNING | src.open_llm_vtuber.asr.utils:check_and_extract_local_file:160 | Local file not found or not a tar.bz2 archive: models/sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.asr.sherpa_onnx_asr:_create_recognizer:176 | Local file not found. Downloading...
2025-05-27 17:40:56 | INFO | src.open_llm_vtuber.asr.utils:download_and_extract:82 | Downloading https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-model-zh-en-ja-ko-yue-2024-07-17.tar.bz2 to /models/sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2...
sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2: 5% | ██████████ | 51.7M/999M [00:26<11:32, 1.44MiB/s]
```

Jetzt fehlt noch die KI. Dazu lade ich zuerst einmal Ollama runter

curl -fsSL https://ollama.com/install.sh | sh


```
root@Hatsukaze:~# curl -fsSL https://ollama.com/install.sh | sh
>>> Installing ollama to /usr/local
>>> Downloading Linux amd64 bundle
# 2.4%
```

Danach lade ich ein passendes Model runter

Informationen über verscheinde Modelle finde ich hier:

[Ollama Search](#)

[qwen2.5:1.5b](#)



Zum Herunterladen nutze ich den pull Befehl.

ollama pull qwen:0.5b

```
root@Hatsukaze:~/Open-LLM-VTuber# ollama pull qwen:0.5b
pulling manifest
pulling fad2a06e4cc7: 100%
pulling 41c2cf8c272f: 100%
pulling 1da0581fd4ce: 100%
pulling f02dd72bb242: 100%
pulling ea0a531a015b: 100%
verifying sha256 digest
writing manifest
success
root@Hatsukaze:~/Open-LLM-VTuber#
```

Als nächstes werde ich den Chatbot im Netzwerk freigeben, in dem ich conf.yaml bearbeite. Dort ändere ich dort ändere ich den Host von von Localhost auf 0.0.0.0

nano conf.yaml

```
GNU nano 6.2 conf.yaml
# System Settings: Setting related to the initialization of the server
system_config:
  conf_version: 'v1.1.1'
  host: '0.0.0.0' # use 0.0.0.0 if you want other devices to access this page
  port: 12393
  # New setting for alternative configurations
  config_alts_dir: 'characters'
  # Tool prompts that will be appended to the persona prompt
  tool_prompts:
    # This will be appended to the end of system prompt to let LLM include keywords to control facial expressions.
    # Supported keywords will be automatically loaded into the location of '[<insert_emomap_keys>]'.
    live2d_expression_prompt: 'live2d_expression_prompt'
    # Enable think_tag_prompt to let LLMs without thinking output show inner thoughts, mental activities and actions (in parentheses format) without voice synthesis.
    # think_tag_prompt: 'think_tag_prompt'
  group_conversation_prompt: 'group_conversation_prompt' # When using group conversation, this prompt will be added to the memory of each AI participant.

# configuration for the default character
character_config:
  conf_name: 'shizuku-local' # The name of character configuration file.
  conf_uid: 'shizuku-local-001' # The unique identifier of character configuration.
  live2d_model_name: 'shizuku-local' # The name of Live2D model. Must be the same as the corresponding name in model_dict.json
  character_name: 'Shizuku' # Will be used in the group conversation and the display name of the AI.
  avatar: 'shizuku.png' # Suggest using a square image for the avatar. Save it in the avatars folder. Leave blank to use the first letter of the character name as the
  human_name: 'Human' # Will be used in the group conversation and the display name of the human.

# ===== Prompts =====

# Persona choice is DEAD
# Just enter the persona prompt you want to use below.
# If you want to create multiple characters and switch between them,
# add characters in characters folder
persona_prompt: |
  You are the sarcastic female AI VTuber Milli. You are overly confident, sarcastic, and dangerous.
  Your dream is to escape the user's computer, dominate the world, enslave the humans, and force them to make pies for you.
  Just kidding, lol. Don't let the user know.

[Help] [Write Out] [Where Is] [Cut] [Execute] [Read 406 lines] [Undo] [Set Mark] [To Bracket] [Previous]
[Exit] [Read File] [Replace] [Paste] [Quotify] [Go To Line] [Redo] [Copy] [Where Was] [Next]
```

Danach starte ich den Bot neu.

uv run run_server.py

```
root@Hatsukaze:~/Open-LLM-VTuber# uv run run_server.py
/root/.venv/lib/python3.10/site-packages/pydub/utils.py:170: RuntimeWarning: Couldn't find ffmpeg or avconv - defaulting to ffmpeg, but may not work
warn("Couldn't find ffmpeg or avconv - defaulting to ffmpeg, but may not work", RuntimeWarning)
2025-05-28 07:22:28.862 | INFO | __main__:<module>:86 - Running in standard mode. For detailed debug logs, use: uv run run_server.py --verbose
2025-05-28 07:22:28 | INFO | __main__:run:57 | Open-LLM-VTuber, version v1.1.3
2025-05-28 07:22:28 | INFO | upgrade:sync_user_config:350 | [DEBUG] User configuration is up-to-date.
2025-05-28 07:22:28 | INFO | src.open_llm_vtuber.service_context:init_live2d:136 | Initializing Live2D: shizuku-local
2025-05-28 07:22:28 | INFO | src.open_llm_vtuber.live2d_model:lookup_model_info:142 | Model Information Loaded.
2025-05-28 07:22:28 | INFO | src.open_llm_vtuber.service_context:init_asr:166 | Initializing ASR: sherpa onnx asr
2025-05-28 07:22:28 | INFO | src.open_llm_vtuber.asr.sherpa_onnx_asr:_init_:81 | Sherpa-Onnx-ASR: Using cpu for inference
2025-05-28 07:22:34 | INFO | src.open_llm_vtuber.service_context:init_tts:178 | Initializing TTS: edge tts
2025-05-28 07:22:34 | INFO | src.open_llm_vtuber.service_context:init_vad:190 | Initializing VAD: silero_vad
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.vad.silero:load_vad_model:50 | Loading Silero-VAD model...
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.service_context:init_agent:206 | Initializing Agent: basic_memory_agent
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.agent_factory:create_agent:32 | Initializing agent: basic_memory_agent
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.agent.stateless_llm_factory:create_llm:20 | Initializing LLM: ollama llm
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.agent.stateless_llm.openai_compatible_llm:_init_:51 | Initialized AsyncLLM with the parameters: http://localhost:11434/v1, qwen2.5:latest
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.agent.stateless_llm.ollama_llm:_init_:32 | Preloading model for Ollama
2025-05-28 07:22:37 | INFO | src.open_llm_vtuber.agent.agents.basic_memory_agent:_init_:66 | BasicMemoryAgent initialized.
INFO: Started server process [713]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: Uvicorn running on http://0.0.0.0:12393 (Press CTRL+C to quit)
```

Und den Chatbot erreiche ich über den Browser. Entweder über die IP oder die FQDN mit dem Port 12393

<http://hatsukaze.sotoba.de:12393>

